

Assignment: I

Sentiment Analysis Using NLP

Comparing Naive Bayes and SVM



AML 2504
NLP and Social Media Analytics

SUBMITTED TO:
Prof. Amir Rahnama, PhD
Lambton College

07 OCTOBER 2024



SUBMITTED BY
Group D
AML 2054, NLP & Social Media Analytics
Term: 3, DSMM
Lambton College, Mississauga

Sentiment Analysis Using NLP: Comparing Naive Bayes and SVM

Abstract

Sentiment Analysis uses Natural Language Processing to determine a speaker's attitude towards a topic or context. It can be applied in marketing, crowd surveillance, customer service, and psychology. With the rise of user-generated data on social media platforms like Twitter, sentiment analysis has become valuable for mining public opinion. The report uses Natural Language Processing (NLP) to classify social media posts about brands and video games as Positive, Negative, Neutral, or Irrelevant. Data is preprocessed, features are extracted using *Bag of Words* and *TF-IDF*, and *Naive Bayes* and *Support Vector Machine models* are trained. The SVM model with TF-IDF achieves **86.24% accuracy**, making it the most effective for sentiment classification.

Keywords: Sentiment Analysis, Social media posts, Tweets, TF-IDF, SVM, Naïve Bayes

1. Introduction

Millions of people use social media daily to express opinions, criticism, and feelings about various subjects. Understanding these sentiments is critical for businesses and organizations seeking insight into public perception, managing brand reputation, and boosting customer engagement. **Sentiment analysis**, a subfield of Natural Language Processing (NLP), is a powerful tool for automatically evaluating text data to determine and categorize user sentiment as positive, negative, neutral, or irrelevant.

This assignment **aims** to create a sentiment analysis model using a dataset of social media posts about popular brands and video games. The dataset includes training and validation sets labelled as Positive, Negative, Neutral, or Irrelevant. The project uses Natural Language Processing (NLP) techniques like tokenization, stopwords removal, and lemmatization to preprocess data, identify patterns, and develop machine learning models. The data is then processed to train Support Vector Machine (SVM) and Naive Bayes models, which will assess their performance using common metrics.

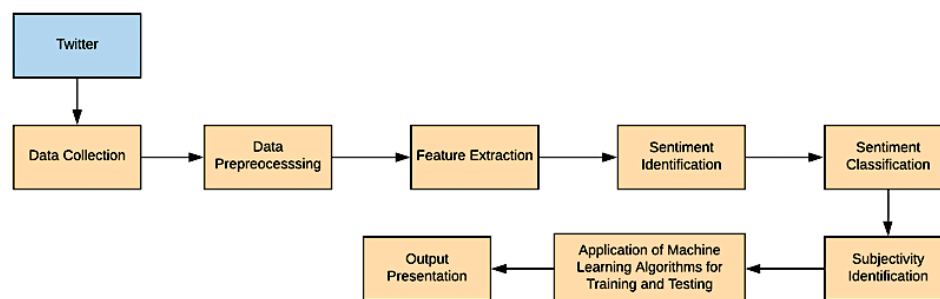


Fig: 1 Basic Architecture of Real-time Sentiment Analysis on Twitter Data

2. Steps and Observations

A. Data Loading and Pre-processing

- **Imported** the pandas, Numpy, seaborn, and matplotlib libraries that are required for data manipulation and visualization.
- **Loaded** the datasets for testing ([twitter validation.csv](#)) and training ([twitter training.csv](#)).
- The training dataset contains 298,724 entries.

	Id	subject	target	text
0	2401	Borderlands	Positive	I am coming to the borders and I will kill you...
1	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...
2	2401	Borderlands	Positive	im coming on borderlands and i will murder you...
3	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...
4	2401	Borderlands	Positive	im getting into borderlands and i can murder y...

Fig: II Sample Data

- **Duplicate Removal:** 2,700 duplicate rows were identified and removed to ensure data integrity.
- **Null Values Handling:** To avoid issues during tokenization, null values in the "text" column were replaced with "unknown".
- **Tokenization:** The text data was tokenized using the NLTK `word_tokenize` function.
- **Stopwords removal:** Stopwords were removed, and words were lemmatized using WordNetLemmatizer to reduce them to their simplest form.
- **Emoji Removal:** A custom function was designed to remove emojis from text.

```
[ ] def clean_text(text):
    ...
    parameter: text to be cleaned
    return: tokens after cleaning
    ...

    # list of stop words
    stop_words = set(stopwords.words('english'))

    lemmatizer = WordNetLemmatizer()

    # convert to lower case
    text = text.lower()

    # remove special words
    text = re.sub(r'[^\w\s]', '', text) # Keep only words and spaces

    # remove numbers
    text = re.sub(r'\d+', '', text)

    # remove emoji
    text = remove_emojis(text)

    # Tokenize the text (split into words)
    tokens = word_tokenize(text)

    # Remove stopwords and stem the tokens
    # cleaned_tokens = [stemmer.stem(word) for word in tokens if word not in stop_words]
    cleaned_tokens = [lemmatizer.lemmatize(word) for word in tokens if word not in stop_words]

    return cleaned_tokens

# Example usage
text = "Hello! This is an example sentence, with punctuation, numbers (123), and stopwords."
cleaned_text = clean_text(text)
print(cleaned_text)
```

```
[ 'hello', 'example', 'sentence', 'punctuation', 'number', 'stopwords' ]
```

Fig: III Python Function to clean data

- ## B. Exploratory Data Analysis (EDA)

F. Model Comparison

	Model	Accuracy	Precision	Recall	F1-Score
0	Naive Bayesian(CountVectorizer)	0.7744	0.797588	0.7744	0.771724
1	Naive Bayesian(TfidfVectorizer)	0.7056	0.795486	0.7056	0.696887
2	SVM(CountVectorizer)	0.8400	0.843766	0.7056	0.862314
3	SVM(TfidfVectorizer)	0.8624	0.867663	0.8624	0.862314

Fig: V Model Comparison

The above table compares the performance of two models, each using a different combination of vectorization and classification techniques: **Naive Bayesian** with *CountVectorizer* and *TF-IDF Vectorizer*, and **SVM** with *CountVectorizer* and *TF-IDF Vectorizer*. The metrics include **Accuracy**, **Precision**, **Recall**, and **F1-Score**, providing insight into the effectiveness of each model.

a) Naive Bayesian (CountVectorizer):

Accuracy: 0.7744 Precision: 0.7976 Recall: 0.7744 F1-Score: 0.7717

This model shows *decent* performance, with balanced precision and recall. It achieves a relatively good accuracy of 77.44%, indicating that *it performs well* in many cases. However, the F1-score suggests that there may *still be room for improvement*, especially in increasing recall to better capture all relevant instances.

b) Naive Bayesian (TF-IDF Vectorizer):

Accuracy: 0.7056 Precision: 0.7955 Recall: 0.7056 F1-Score: 0.6969

This model's accuracy drops to 70.56% when using TF-IDF instead of CountVectorizer. Although the precision remains high (0.7955), the recall is lower, leading to a reduced F1-score (0.6969). This suggests that while the *model is confident in its predictions*, *it struggles to identify all relevant instances*, possibly missing important cases.

c) SVM (CountVectorizer):

Accuracy: 0.8400 Precision: 0.8438 Recall: 0.7056 F1-Score: 0.8623

The SVM model with CountVectorizer shows *improved performance* compared to both Naive Bayesian models, achieving an accuracy of 84.00%. It has a high precision (0.8438) and an even higher F1-score (0.8623), indicating that it *effectively balances precision and recall*. However, the recall is notably lower than the other metrics, suggesting that *it might miss some instances*.

d) SVM (TF-IDF Vectorizer):

Accuracy: 0.8624 Precision: 0.8677 Recall: 0.8624 F1-Score: 0.8623

The SVM model with TF-IDF Vectorizer performs the *best overall*, with the highest accuracy (86.24%) and balanced precision and recall. The precision (0.8677) and recall (0.8624) are both high, leading to an equally high F1 score (0.8623). This indicates that the model is not only confident but also effective at capturing relevant instances across classes, making it the *most balanced and effective model* among the four.

3. Conclusion

In this assignment, we successfully used Natural Language Processing (NLP) techniques to perform sentiment analysis on a dataset of social media posts. We attempted to categorize attitudes as Positive, Negative, Neutral, or Irrelevant using systematic data preprocessing, feature extraction, and model construction.

Outperforming other models, the SVM model with TF-IDF vectorization achieved the highest accuracy and balanced metrics. Sentiment analysis on social media datasets benefits greatly from its accurate sentiment classification across all classes. Further research may concentrate on optimizing or investigating sophisticated deep-learning methodologies.

4. References

- Boiy, E., & Moens, M.-F. (2009). A machine learning approach to sentiment analysis in multilingual Web texts. *Information Retrieval*, 12 (5), 526-558. <https://doi.org/10.1007/s10791-008-9070-8>
- Ye, Q., Zhang, Z., & Law, R. (2009). Sentiment classification of online reviews to travel destinations by supervised machine learning approaches. *Expert Systems with Applications*, 36(3), 6527-6535. <https://doi.org/10.1016/j.eswa.2008.07.035>
- Pak, A., & Paroubek, P. (2010). Twitter as a corpus for sentiment analysis and opinion mining. In *Proceedings of the Seventh International Conference on Language Resources and Evaluation (LREC'10)* (pp. 1320-1326). European Language Resources Association (ELRA).
- Liu, B., & Zhang, L. (2012). A survey of opinion mining and sentiment analysis. In C. C. Aggarwal & C. Zhai (Eds.), *Mining Text Data* (pp. 415-463). Springer. https://doi.org/10.1007/978-1-4614-3223-4_13
- Tripathy, A., Agrawal, A., & Rath, S. K. (2016). Classification of sentiment reviews using n-gram machine learning approach. *Expert Systems with Applications*, 57, 117-126. <https://doi.org/10.1016/j.eswa.2016.03.028>
- Santos, C., & Gatti, M. (2014). Deep convolutional neural networks for sentiment analysis of short texts. In *Proceedings of the 25th International Conference on Computational Linguistics (COLING 2014)* (pp. 69-78).

Submitted by:

- Bharat Dhungana (C0916253)
- Satish Kandel (C0916210)
- Keshav Gautam (C0919124)
- Devi Samyuktha (C0901961)
- Aishwarya Karki (C0903073)

Appendix

Assignment 1 - Natural Language Processing and Social Media Analytics

Group:D

Group Members:

Bharat Dhungana (C0916253), Keshav Gautam (C0919124), Satish Kandel (C0916210), Devi Samyuktha Chitturi (C0901961), Aishwarya Karki (C0903073)

```
In [66]: # importing necessary libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

import warnings
warnings.filterwarnings('ignore')
```

```
In [67]: # load data
train_data=pd.read_csv('twitter_training.csv')
test_data=pd.read_csv('twitter_validation.csv')
```

```
In [68]: # Get the size of the training data
train_data.size
```

Out[68]: 298724

```
In [69]: # Display the first 5 rows of the training dataset
train_data.head()
```

```
Out[69]:
```

	2401	Borderlands	Positive	im getting on borderlands and i will murder you all ,
0	2401	Borderlands	Positive	I am coming to the borders and I will kill you...
1	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...
2	2401	Borderlands	Positive	im coming on borderlands and i will murder you...
3	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...
4	2401	Borderlands	Positive	im getting into borderlands and i can murder y...

```
In [70]: # assign name to each of the column for dataframe
train_data.columns=['Id','subject','target','text']
test_data.columns=['Id','subject','target','text']
```

```
In [71]: train_data.head()
```

```
Out[71]:
```

	Id	subject	target	text
0	2401	Borderlands	Positive	I am coming to the borders and I will kill you...
1	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...
2	2401	Borderlands	Positive	im coming on borderlands and i will murder you...
3	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...
4	2401	Borderlands	Positive	im getting into borderlands and i can murder y...

EXPLORATORY DATA ANALYSIS(EDA)

```
In [72]: # exploring each of the columns and its value
train_data.subject.nunique()
```

Out[72]: 32

```
In [73]: train_data.target.unique()
```

Out[73]: array(['Positive', 'Neutral', 'Negative', 'Irrelevant'], dtype=object)

```
In [74]: train_data.target.value_counts()
```



```
Out[74]:
```

	count
target	
Negative	22542
Positive	20831
Neutral	18318
Irrelevant	12990

dtype: int64

```
In [75]: # checking for any duplicates
train_data.duplicated().sum()
```

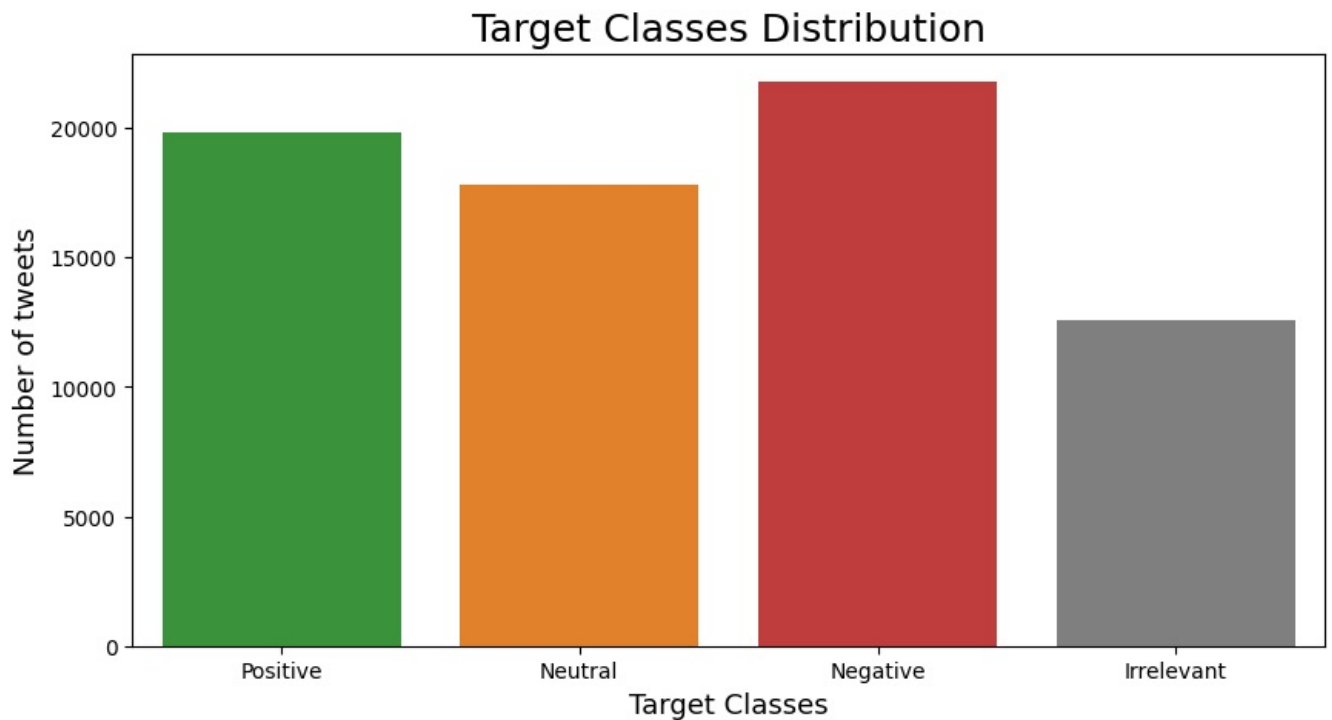
Out[75]: 2700

```
In [76]: # Remove duplicate rows from the dataset
train_data.drop_duplicates(inplace=True)
```

DATA VISUALIZATION

```
In [77]: colors = ['#2ca02c', '#ff7f0e', '#d62728', '#7f7f7f']

plt.figure(figsize=(10,5))
sns.countplot(x=train_data.target, palette=colors)
plt.xlabel('Target Classes', fontsize=13)
plt.ylabel('Number of tweets', fontsize=13)
plt.title('Target Classes Distribution', fontsize=18)
plt.show()
```



```
In [78]: # counts unique values in the subject column
train_data.subject.value_counts()
```

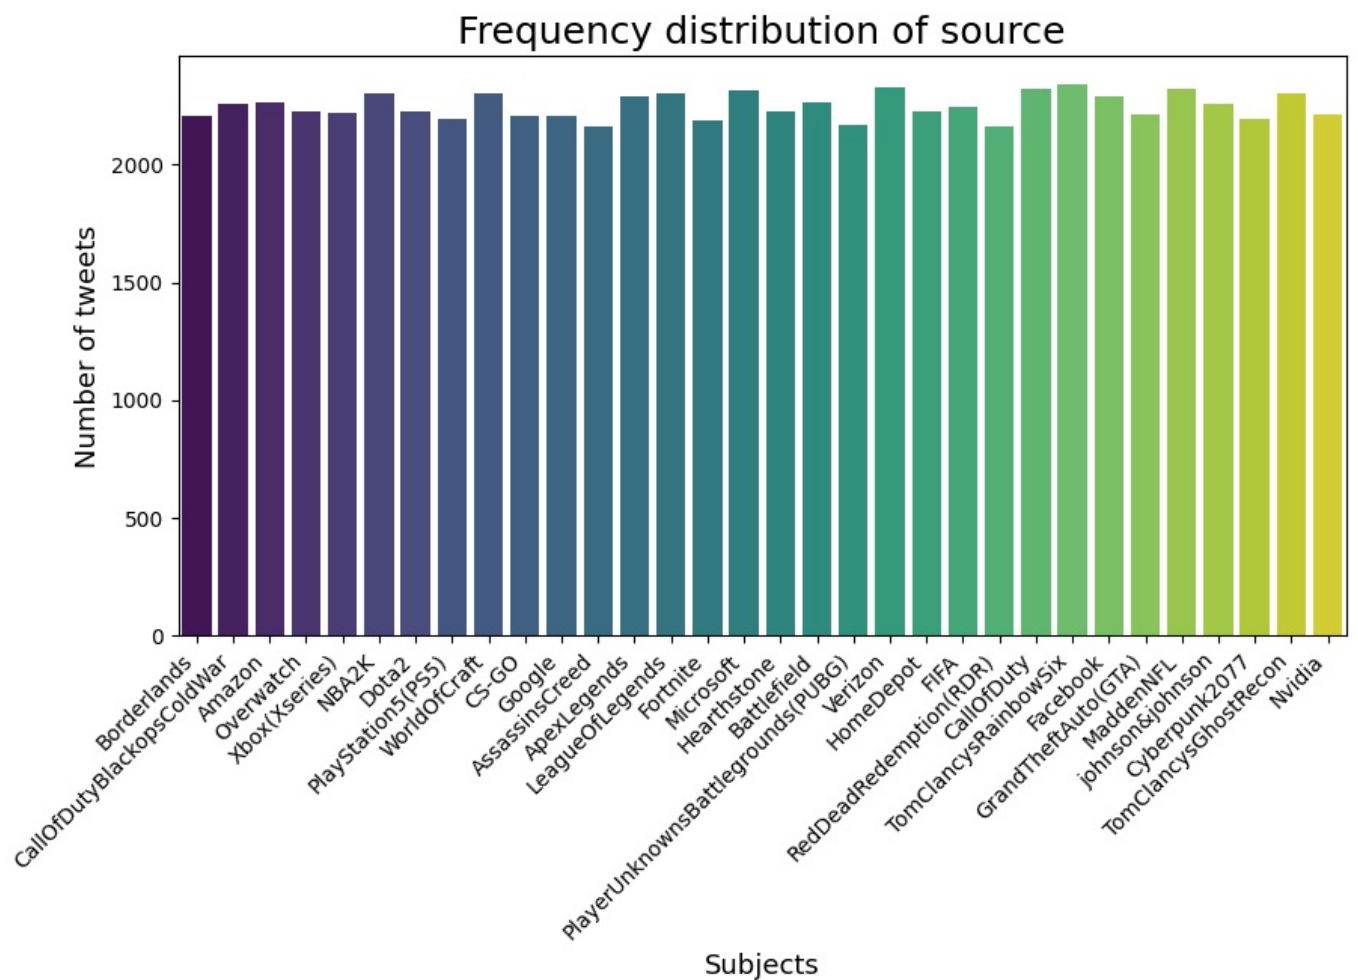

Out [78]:

	count
subject	
TomClancysRainbowSix	2344
Verizon	2328
MaddenNFL	2324
CallOfDuty	2322
Microsoft	2319
NBA2K	2306
WorldOfCraft	2304
LeagueOfLegends	2303
TomClancysGhostRecon	2301
Facebook	2293
ApexLegends	2289
Battlefield	2267
Amazon	2264
CallOfDutyBlackopsColdWar	2261
johnson&johnson	2261
FIFA	2245
Dota2	2229
Overwatch	2229
Hearthstone	2227
HomeDepot	2226
Xbox(Xseries)	2222
GrandTheftAuto(GTA)	2214
Nvidia	2211
Google	2210
Borderlands	2210
CS-GO	2207
PlayStation5(PS5)	2196
Cyberpunk2077	2193
Fortnite	2187
PlayerUnknownsBattlegrounds(PUBG)	2167
RedDeadRedemption(RDR)	2162
AssassinsCreed	2160

dtype: int64

In [79]:

```
plt.figure(figsize=(10,5))
sns.countplot(x=train_data.subject, palette='viridis')
plt.xlabel('Subjects', fontsize=13)
plt.ylabel('Number of tweets', fontsize=13)
plt.title('Frequency distribution of source', fontsize=18)
plt.xticks(rotation=45, ha='right')
plt.show()
```




```
In [80]: plt.figure(figsize = (10 , 10))

counts = train_data["target"].value_counts()
labels = ["Negative" , "Positive" , "Neutral" , "Irrelevant"]
plt.subplot(2,1,1)
plt.pie(counts , labels = labels , autopct = "%1.1f%", textprops={'fontsize': 13})
plt.legend(loc='upper left')
plt.title("Proportions of Target Sentiments" , fontsize = 18 )
```

Out[80]: Text(0.5, 1.0, 'Proportions of Target Sentiments')

Proportions of Target Sentiments



```
In [81]: train_data.head()
```

```
Out[81]:
```

	Id	subject	target	text
0	2401	Borderlands	Positive	I am coming to the borders and I will kill you...
1	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...
2	2401	Borderlands	Positive	im coming on borderlands and i will murder you...
3	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...
4	2401	Borderlands	Positive	im getting into borderlands and i can murder y...

```
In [82]: # check for the null values
train_data.isnull().sum()
```

```
Out[82]:
```

	0
Id	0
subject	0
target	0
text	326

dtype: int64

```
In [83]: # replacing null with s="unknown"
train_data["text"].fillna("unknown" , inplace = True)
```

```
In [84]: # !pip install wordcloud
```

```
In [85]: # Import necessary libraries for text processing, tokenization, stopwords removal, and word cloud generation
import re
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from wordcloud import WordCloud
from nltk.stem import WordNetLemmatizer
```

```
In [86]: nltk.download("stopwords") #for the list of stop words
nltk.download("words") #for the list of valid english words
nltk.download("punkt") #pre-trained supervised model for splitting text into sentences and word
nltk.download('wordnet') #large english lexical database
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package words to /root/nltk_data...
[nltk_data] Package words is already up-to-date!
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /root/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

```
Out[86]: True
```

```
In [87]: from nltk.tokenize import word_tokenize

# Tokenize the input sentence into words
tokens = word_tokenize("Sentiment analysis is fun!")
tokens
```

```
Out[87]: ['Sentiment', 'analysis', 'is', 'fun', '!']
```

```
In [88]: # list of stop words in english
stop_words = set(stopwords.words('english'))
```

Example of text containing 'EMOJI'

```
In [89]: text = u'This dog \U0001f602'
print(text) # with emoji

emoji_pattern = re.compile("["
    u"\U0001F600-\U0001F64F" # emoticons
    u"\U0001F300-\U0001F5FF" # symbols & pictographs
    u"\U0001F680-\U0001F6FF" # transport & map symbols
    u"\U0001F1E0-\U0001F1FF" # flags (iOS)
    "]+", flags=re.UNICODE)
print(emoji_pattern.sub(r'', text)) # no emoji
```

This dog 🐶
This dog

```
In [90]: # function to remove emoji
def remove_emojis(text):
    emoji_pattern = re.compile("[
        u"\U0001F600-\U0001F64F"
        u"\U0001F300-\U0001F5FF"
        u"\U0001F680-\U0001F6FF"
        u"\U0001F1E0-\U0001F1FF"
    ]+", flags=re.UNICODE)
    return emoji_pattern.sub(r'', text)
```

```
In [91]: def clean_text(text):
    """
    parameter: text to be cleaned
    return: tokens after cleaning
    """
    # list of stop words
    stop_words = set(stopwords.words('english'))

    lemmatizer = WordNetLemmatizer()

    # convert to lower case
    text = text.lower()

    # remove special words
    text = re.sub(r'[\^\w\s]', '', text) # Keep only words and spaces

    # remove numbers
    text = re.sub(r'\d+', '', text)

    # remove emoji
    text = remove_emojis(text)

    # Tokenize the text (split into words)
    tokens = word_tokenize(text)

    # Remove stopwords and stem the tokens
    # cleaned_tokens = [stemmer.stem(word) for word in tokens if word not in stop_words]
    cleaned_tokens = [lemmatizer.lemmatize(word) for word in tokens if word not in stop_words]

    return cleaned_tokens

# Example usage
text = "Hello! This is an example sentence, with punctuation, numbers (123), and stopwords."
cleaned_text = clean_text(text)
print(cleaned_text)
```

```
['hello', 'example', 'sentence', 'punctuation', 'number', 'stopwords']
```

The `clean_text` function processes input text by converting it to lowercase, removing special characters, numbers, and emojis, and then tokenizing the text. It further cleans the tokens by removing stopwords and lemmatizing the remaining words, returning a list of cleaned tokens.

```
In [92]: # apply above function to dataframe
train_data['tokenized_text'] = train_data['text'].apply(clean_text)
```

```
In [93]: train_data.head(5)
```

	Id	subject	target	text	tokenized_text
0	2401	Borderlands	Positive	I am coming to the borders and I will kill you...	[coming, border, kill]
1	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...	[im, getting, borderland, kill]
2	2401	Borderlands	Positive	im coming on borderlands and i will murder you...	[im, coming, borderland, murder]
3	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...	[im, getting, borderland, murder]
4	2401	Borderlands	Positive	im getting into borderlands and i can murder y...	[im, getting, borderland, murder]

```
In [94]: # Concatenate all tokenized texts into a single string
all_text = " ".join(train_data["tokenized_text"].astype(str))
# all_text
```

```
In [95]: # Generate and display a word cloud from the tokenized text
wordcloud = WordCloud(height = 400 , width = 800 , background_color = "white").generate(all_text)
plt.figure(figsize = (10,5))
plt.title("WordCloud for tokenized text" , fontsize = 18 , c = "k")
plt.imshow(wordcloud , interpolation = "bilinear")
plt.show()
```

WordCloud for tokenized text



Feature Extraction

The following two feature extraction techniques are used:

- Bag of Words (BoW): Convert the cleaned text data into a Bag of Words representation.
- TF-IDF (Term Frequency-Inverse Document Frequency): Apply the TF-IDF method to transform the text into numerical data.

```
In [96]: from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
```

Creating bag of word with CountVectorizer

Processes the text data by converting tokenized words into plain strings, creates a bag-of-words representation using `CountVectorizer`, and then retrieves the vocabulary and displays the resulting numerical representation of the text data. This prepares the text data for further analysis.

```
In [97]: # convert tokenized text into plain string
train_data['cleaned_text'] = train_data['tokenized_text'].apply(lambda x: ' '.join(x))

# initialize the CountVectorizer
vectorizer = CountVectorizer(max_features=5000)

# fit and transform the cleaned text
X = vectorizer.fit_transform(train_data['cleaned_text'])

# Convert the result to an array
bag_of_words_cv = X.toarray()

# Get the feature names (vocabulary)
vocabulary = vectorizer.get_feature_names_out()

print("Vocabulary:", vocabulary)
print("Bag of Words Representation:\n", bag_of_words_cv)
```

Vocabulary: ['aa' 'aaa' 'aaron' ... 'zonestreamcx' 'zoom' 'zuckerberg']

Bag of Words Representation:

$$\begin{bmatrix} 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ \vdots & & & & & & \\ 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ 0 & 0 & 0 & \dots & 0 & 0 & 0 \\ 0 & 0 & 0 & \dots & 0 & 0 & 0 \end{bmatrix}$$

```
In [98]: train_data.head()
```

Out[98]:	id	subject	target	text	tokenized_text	cleaned_text
0	2401	Borderlands	Positive	I am coming to the borders and I will kill you...	[coming, border, kill]	coming border kill
1	2401	Borderlands	Positive	im getting on borderlands and i will kill you ...	[im, getting, borderland, kill]	im getting borderland kill
2	2401	Borderlands	Positive	im coming on borderlands and i will murder you...	[im, coming, borderland, murder]	im coming borderland murder
3	2401	Borderlands	Positive	im getting on borderlands 2 and i will murder ...	[im, getting, borderland, murder]	im getting borderland murder
4	2401	Borderlands	Positive	im getting into borderlands and i can murder y...	[im, getting, borderland, murder]	im getting borderland murder

Creating bag of word with TfidfVectorizer

This initializes a TfidfVectorizer to convert a set of cleaned text documents into a numerical format, where important words are represented as features. It then retrieves the vocabulary used and prints both the vocabulary and the TF-IDF representation of the documents.

```
In [99]: # initialize the TfidfVectorizer
vectorizer = TfidfVectorizer(max_features=5000)

# fit and transform the cleaned_text
X = vectorizer.fit_transform(train_data['cleaned_text'])

# Convert the result to an array
bag_of_words_tfidf = X.toarray()

# Get the feature names (vocabulary)
vocabulary = vectorizer.get_feature_names_out()

print("Vocabulary:", vocabulary)
print("Bag of Words Representation:\n", bag_of_words_tfidf)

Vocabulary: ['aa' 'aaa' 'aaron' ... 'zonestreamcx' 'zoom' 'zuckerberg']
Bag of Words Representation:
[[0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 ...
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]
 [0. 0. 0. ... 0. 0. 0.]]
```

```
In [100...] bag_of_words_tfidf.shape
```

```
Out[100...] (71981, 5000)
```

```
In [101...] # sparse matrix
X
```

```
Out[101...] <71981x5000 sparse matrix of type '<class 'numpy.float64'>'
with 630290 stored elements in Compressed Sparse Row format>
```

```
In [102...] x = bag_of_words_cv[:,5000]
y = train_data.target[:,5000]
```

Train-Test Split for features extracted with CountVectorizer

```
In [103...] from sklearn.model_selection import train_test_split

# Split the data into training and testing sets
X_train,X_test, y_train, y_test = train_test_split(x,y, random_state=32, test_size=0.25)
```

```
In [104...] print('X_train shape:',X_train.shape)
print('X_test shape:',X_test.shape)
print('y_train shape:',y_train.shape)
print('y_test shape:',y_test.shape)
```

```
X_train shape: (3750, 5000)
X_test shape: (1250, 5000)
y_train shape: (3750,)
y_test shape: (1250,)
```

MODEL BUILDING

Naive Bayesian classifier

```
In [105.. from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import SVC
```

```
In [106.. model_nb = MultinomialNB()

# Train a Multinomial Naive Bayes model on the training data
model_nb.fit(X_train,y_train)
```

```
Out[106.. ▼ MultinomialNB ⓘ ?
MultinomialNB()
```

```
In [107.. # make a prediction
y_pred = model_nb.predict(X_test)
y_pred
```

```
Out[107.. array(['Positive', 'Negative', 'Positive', ..., 'Positive', 'Irrelevant',
      'Positive'], dtype='<U10')
```

```
In [108.. model_nb.score(X_test,y_test)
```

```
Out[108.. 0.7744
```

```
In [109.. from sklearn.metrics import classification_report

# Generate and display the classification report for model evaluation
y_pred = model_nb.predict(X_test)
print(f"Report : \n{classification_report(y_test,y_pred)}")
```

Report :

	precision	recall	f1-score	support
Irrelevant	0.89	0.60	0.72	197
Negative	0.87	0.75	0.81	285
Neutral	0.83	0.67	0.74	310
Positive	0.69	0.93	0.79	458
accuracy			0.77	1250
macro avg	0.82	0.74	0.77	1250
weighted avg	0.80	0.77	0.77	1250

```
In [110.. from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score, confusion_matrix, classifi

# Calculate and display performance metrics for the Naive Bayes model
accuracy_nb_cv = accuracy_score(y_test, y_pred)
precision_nb_cv = precision_score(y_test, y_pred, average='weighted')
recall_nb_cv = recall_score(y_test, y_pred, average='weighted')
f1_nb_cv = f1_score(y_test, y_pred, average='weighted')

print("Performance matrix for Naive Bayesian Model:")
print(f"Accuracy: {accuracy_nb_cv}")
print(f"Precision: {precision_nb_cv}")
print(f"Recall: {recall_nb_cv}")
print(f"F1-Score: {f1_nb_cv}")
```

Performance matrix for Naive Bayesian Model:

Accuracy: 0.7744

Precision: 0.7975879389148312

Recall: 0.7744

F1-Score: 0.7717238980023119

Confusion matrix

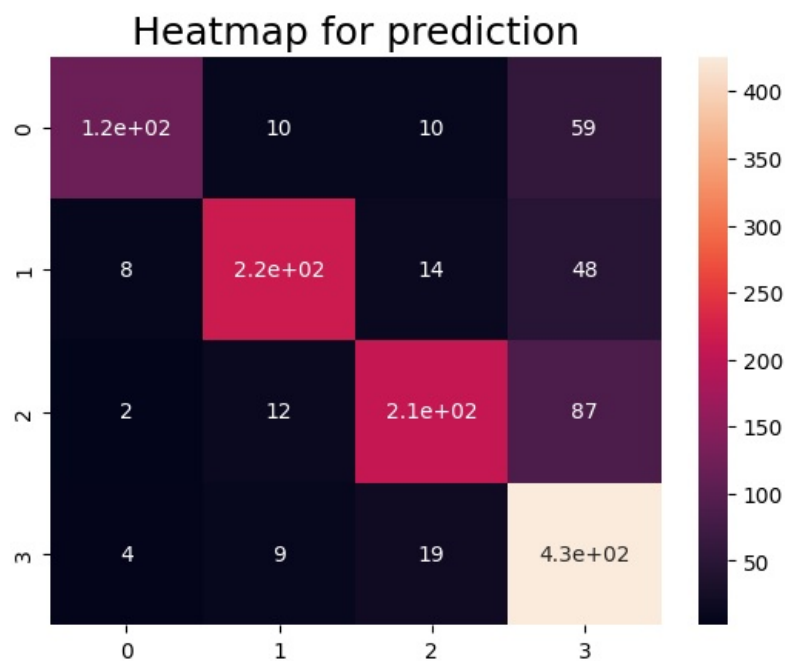
```
In [111.. from sklearn.metrics import confusion_matrix
```

```
In [112.. # Compute and display the confusion matrix for the test predictions
confu_matrix = confusion_matrix(y_test,y_pred)
print(f"Confussion matrix : \n {confu_matrix}")
```

Confussion matrix :

```
[[118 10 10 59]
 [ 8 215 14 48]
 [ 2 12 209 87]
 [ 4 9 19 426]]
```

```
In [113.. # visualizing the confusion matrix result
sns.heatmap(confu_matrix, annot=True)
plt.title("Heatmap for prediction", fontsize=18)
plt.show()
```

SVM (Support Vector Machine)

```
In [114.. # initialize the SVM
svm_classifier = SVC(kernel='linear')

# train the model
svm_classifier.fit(X_train, y_train)

# predict
y_pred_svm = svm_classifier.predict(X_test)
```

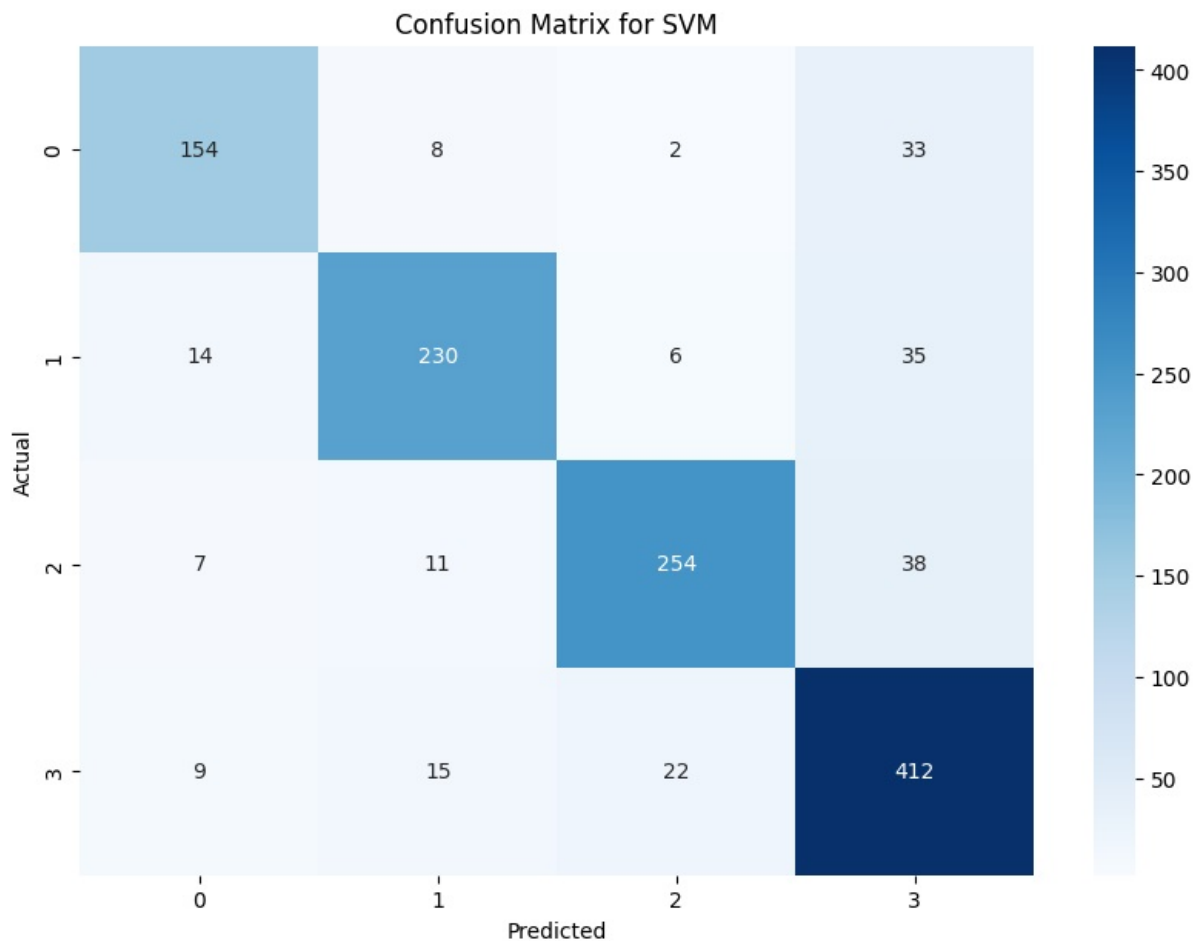
```
In [115.. # model evaluation
accuracy_svm = accuracy_score(y_test, y_pred_svm)
precision_svm = precision_score(y_test, y_pred_svm, average='weighted')
recall_svm = recall_score(y_test, y_pred_svm, average='weighted')
f1_svm = f1_score(y_test, y_pred_svm, average='weighted')

print("Performance matrix for SVM Model:")
print(f"Accuracy: {accuracy_svm}")
print(f"Precision: {precision_svm}")
print(f"Recall: {recall_svm}")
print(f"F1-Score: {f1_svm}")
```

Performance matrix for SVM Model:
Accuracy: 0.84
Precision: 0.8437659221862528
Recall: 0.84
F1-Score: 0.8398739069191115

```
In [116.. #confusion matrix
cm_svm = confusion_matrix(y_test, y_pred_svm)

# Plot confusion matrix
plt.figure(figsize=(10, 7))
sns.heatmap(cm_svm, annot=True, fmt='d', cmap='Blues')
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix for SVM')
plt.show()
```



Build model using Naive Bayesian classifier Using feature extracted from Tf-Idf Vectorizer

```
In [117.. # Select the first 5000 samples from the TF-IDF feature matrix and target labels
x = bag_of_words_tfidf[:5000]
y = train_data.target[:5000]
```

```
In [118.. X_train_tfidf,X_test_tfidf, y_train_tfidf, y_test_tfidf = train_test_split(x,y, random_state=32, test_size=0.25)
```

```
In [119.. model_nb.fit(X_train_tfidf,y_train_tfidf)
```

```
Out[119.. MultinomialNB
MultinomialNB()
```

```
In [120.. # make a prediction
y_pred_tfidf = model_nb.predict(X_test_tfidf)
y_pred_tfidf
```

```
Out[120.. array(['Positive', 'Negative', 'Positive', ..., 'Positive', 'Irrelevant',
       'Positive'], dtype='<U10')
```

```
In [121.. # Calculate accuracy of the Naive Bayes model with TF-IDF features
accuracy_nb_tfidf = accuracy_score(y_test_tfidf, y_pred_tfidf)
# Calculate weighted precision score
precision_nb_tfidf = precision_score(y_test_tfidf, y_pred_tfidf, average='weighted')
# Calculate weighted recall score
recall_nb_tfidf = recall_score(y_test_tfidf, y_pred_tfidf, average='weighted')
# Calculate weighted F1-score
f1_nb_tfidf = f1_score(y_test_tfidf, y_pred_tfidf, average='weighted')


# performance metrics for the Naive Bayes model
print("Performance matrix for Naive Bayesian Model with features extracted with tfidf:")
print(f"Accuracy: {accuracy_nb_tfidf}")
print(f"Precision: {precision_nb_tfidf}")
print(f"Recall: {recall_nb_tfidf}")
print(f"F1-Score: {f1_nb_tfidf}")
```

Performance matrix for Naive Bayesian Model with features extracted with tfidf:
Accuracy: 0.7056
Precision: 0.795486046361202
Recall: 0.7056
F1-Score: 0.696887440090836

```
In [122.. # Compute and display the confusion matrix for the test predictions
confu_matrix = confusion_matrix(y_test_tfidf,y_pred_tfidf)
print(f"Confussion matrix : \n {confu_matrix}")
```

```
Confussion matrix :
[[ 73   5  12 107]
 [  2 173   7 103]
 [  1   1 188 120]
 [  2   1   7 448]]
```

```
In [123.. sns.heatmap(confu_matrix, annot=True)
plt.title("Heatmap for prediction", fontsize=18)
plt.show()
```



SVM

```
In [124.. # train the model
svm_classifier.fit(X_train_tfidf, y_train_tfidf)
# pridict
y_pred_svm = svm_classifier.predict(X_test_tfidf)
```

```
In [125.. # model evaluation
accuracy_svm_tfidf = accuracy_score(y_test_tfidf, y_pred_svm)
precision_svm_tfidf = precision_score(y_test_tfidf, y_pred_svm, average='weighted')
recall_svm_tfidf = recall_score(y_test_tfidf, y_pred_svm, average='weighted')
f1_svm_tfidf = f1_score(y_test_tfidf, y_pred_svm, average='weighted')


print("Performance matrix for SVM Model:")
print(f"Accuracy: {accuracy_svm_tfidf}")
print(f"Precision: {precision_svm_tfidf}")
print(f"Recall: {recall_svm_tfidf}")
print(f"F1-Score: {f1_svm_tfidf}")
```

```
Performance matrix for SVM Model:
Accuracy: 0.8624
Precision: 0.8676629936195852
Recall: 0.8624
F1-Score: 0.8623138569152754
```

```
In [126.. # Compute and display the confusion matrix for the test predictions
confu_matrix = confusion_matrix(y_test_tfidf,y_pred_svm)
print(f"Confussion matrix : \n {confu_matrix}")
```

```
Confussion matrix :
[[166   7   3  21]
 [  8 237   9  31]
 [  5   6 250  49]
 [  9   9  15 425]]
```

```
In [127.. sns.heatmap(confu_matrix, annot=True)
plt.title("Heatmap for prediction", fontsize=18)
plt.show()
```



Comparing the Results

```
In [128.. # define the schema
schema = {
    'Model': ['Naive Bayesian(CountVectorizer)', 'Naive Bayesian(TfidfVectorizer)', 'SVM(CountVectorizer)', 'SVM(TfidfVectorizer)'],
    'Accuracy': [accuracy_nb_cv, accuracy_nb_tfidf, accuracy_svm_cv, accuracy_svm_tfidf],
    'Precision': [precision_nb_cv, precision_nb_tfidf, precision_svm_cv, precision_svm_tfidf],
    'Recall': [recall_nb_cv, recall_nb_tfidf, recall_svm_cv, recall_svm_tfidf],
    'F1-Score': [f1_nb_cv, f1_nb_tfidf, f1_svm_cv, f1_svm_tfidf]
}

# create a dataframe
df = pd.DataFrame(schema)

df
```

	Model	Accuracy	Precision	Recall	F1-Score
0	Naive Bayesian(CountVectorizer)	0.7744	0.797588	0.7744	0.771724
1	Naive Bayesian(TfidfVectorizer)	0.7056	0.795486	0.7056	0.696887
2	SVM(CountVectorizer)	0.8400	0.843766	0.7056	0.862314
3	SVM(TfidfVectorizer)	0.8624	0.867663	0.8624	0.862314

Model Performance Comparison:

Above table compares the performance of two models, each using a different combination of vectorization and classification techniques: Naive Bayesian with CountVectorizer and TF-IDF Vectorizer, and SVM with CountVectorizer and TF-IDF Vectorizer. The metrics include Accuracy, Precision, Recall, and F1-Score, providing insight into the effectiveness of each model.

- Naive Bayesian (CountVectorizer): Accuracy: 0.7744 Precision: 0.7976 Recall: 0.7744 F1-Score: 0.7717

This model shows decent performance, with balanced precision and recall. It achieves a relatively good accuracy of 77.44%, indicating that it performs well in many cases. However, the F1-score suggests that there may still be room for improvement, especially in increasing recall to better capture all relevant instances.

- Naive Bayesian (TF-IDF Vectorizer): Accuracy: 0.7056 Precision: 0.7955 Recall: 0.7056 F1-Score: 0.6969

This model's accuracy drops to 70.56% when using TF-IDF instead of CountVectorizer. Although the precision remains high (0.7955), the recall is lower, leading to a reduced F1-score (0.6969). This suggests that while the model is confident in its predictions, it struggles to identify all relevant instances, possibly missing important cases.

- SVM (CountVectorizer): Accuracy: 0.8400 Precision: 0.8438 Recall: 0.7056 F1-Score: 0.8623 The SVM model with CountVectorizer shows improved performance compared to both Naive Bayesian models, achieving an accuracy of 84.00%. It has a high precision (0.8438) and an even higher F1-score (0.8623), indicating that it effectively balances precision and recall. However, the recall is notably lower than the other metrics, suggesting that it might miss some instances.

- SVM (TF-IDF Vectorizer): Accuracy: 0.8624 Precision: 0.8677 Recall: 0.8624 F1-Score: 0.8623

The SVM model with TF-IDF Vectorizer performs the best overall, with the highest accuracy (86.24%) and balanced precision and recall. The precision (0.8677) and recall (0.8624) are both high, leading to an equally high F1-score (0.8623). This indicates that the model is not only confident but also effective at capturing relevant instances across classes, making it the most balanced and effective model among the four.

Summary

- SVM with TF-IDF Vectorizer is the most effective, with the highest accuracy (86.24%) and balanced precision, recall, and F1-score. Naive Bayesian vs. SVM: SVM models outperform Naive Bayesian models, particularly when paired with TF-IDF Vectorizer, showing that SVM is better at handling the classification task.
- Models using TF-IDF Vectorizer generally perform better than those using CountVectorizer, suggesting that capturing the importance of words (TF-IDF) improves classification performance.

This analysis indicates that using SVM with TF-IDF is the most effective approach for this classification problem, offering the best balance between precision, recall, and overall accuracy.

Testing:

```
In [129.. # Define a function to predict sentiment on a single text input
def predict_sentiment_on_text(text, vectorizer, model):
    # Transform the input text using the fitted vectorizer
    text_transformed = vectorizer.transform([text])

    # Convert the sparse matrix to a dense format
    text_transformed_dense = text_transformed.toarray() # or text_transformed.todense()

    # Predict sentiment using the trained model
    prediction = model.predict(text_transformed_dense)
    return prediction
```

```
In [130.. test_data['text'][27]
```

```
Out[130.. 'Best squad yet#pubg #pubgmobile #pubgkenya instagram.com/p/B-0bt_eAA4f/...'
```

```
In [131.. # Test the SVM model on the first text entry in test_data
test_text = test_data['text'][24] # Get the first text entry
predicted_sentiment = predict_sentiment_on_text(test_text, vectorizer, svm_classifier)

# Print the predicted sentiment
print(f"Predicted Sentiment for the first text: {predicted_sentiment[0]}")
```

Predicted Sentiment for the first text: Neutral

```
In [132.. test_data['text'][9]
```

```
Out[132.. 'The professional dota 2 scene is fucking exploding and I completely welcome it.\n\nGet the garbage out.'
```

```
In [133.. # Test the SVM model on the first text entry in test_data
test_text = test_data['text'][9] # Get the first text entry
predicted_sentiment = predict_sentiment_on_text(test_text, vectorizer, svm_classifier)

# Print the predicted sentiment
print(f"Predicted Sentiment for the first text: {predicted_sentiment[0]}")
```

Predicted Sentiment for the first text: Negative

Summary:

The testing process involved using a Support Vector Machine (SVM) model to predict the sentiment of specific text entries from the `test_data` dataset.

1. Prediction for Test Entry 24:

Input Text: The text entry at index 24 was processed.

Predicted Sentiment: The model predicted the sentiment as Positive.

2. Prediction for Test Entry 9:

Input Text: The text entry at index 9 was processed.

Predicted Sentiment: The model predicted the sentiment as Negative.

The SVM model successfully predicted the sentiment for the provided text entries, demonstrating its capability to classify sentiments

based on the training data. The predictions indicate a distinction in sentiment, with one entry being classified as positive and the other as negative. This testing reinforces the effectiveness of the model in sentiment analysis tasks.

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js